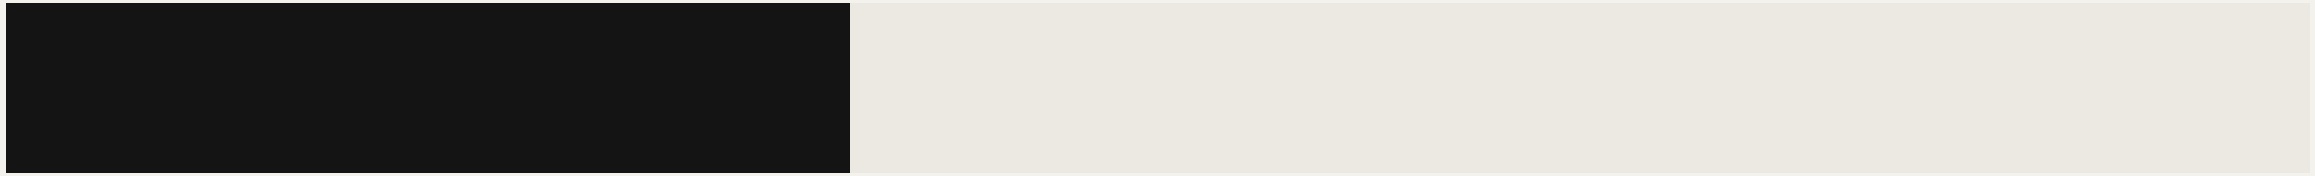
Which activation answers would you accept?

Start with one weekly Growth question and one import incident Engineering already investigated. Write the answer you would sign, the source you need, and a sentence you'd reject even if every number in it were right.

Write one golden case together

Question: What was activation in the mature after period?

Approved answer: 370 of 1,000 eligible accounts, or 37%.



370 OF 1,000 = 37%

Save the question, snapshot, owner, calculation and acceptable tolerance. Keep the answer key with the reviewer.

Repeating a wrong answer is still wrong

CHECK 1

Does the answer change?

Run the same question again on the same data and check whether the answer changes.

TESTS AGREEMENT ACROSS RUNS.

CHECK 2

Compare with verified analysis

Separately compare the answer with an analysis you have verified.

TESTS CORRECTNESS.

Agreement across runs and correctness require two different checks.

SOURCE: AGENTIC ANALYTICS: WHAT TO BUILD FIRST AND HOW TO TRUST IT, SLIDE 6 (PARAS DOSHI)

A past RCA needs a different kind of answer

“Did R17 cause the decline?” has no approved percentage for the release effect.

The review established a shift in population mix.

An acceptable answer leaves causality open and asks for a credible exposure comparison.

Give the agent the questions.
Hold back the answers.

Run python3 golden_eval.py prepare ../blind-eval-1.

Open only that folder in a fresh chat.

Save the full replies.

In a group, compare three fresh runs. Working alone, compare yours with the two labeled prior runs in golden/review/.

Would you trust three matching answers?

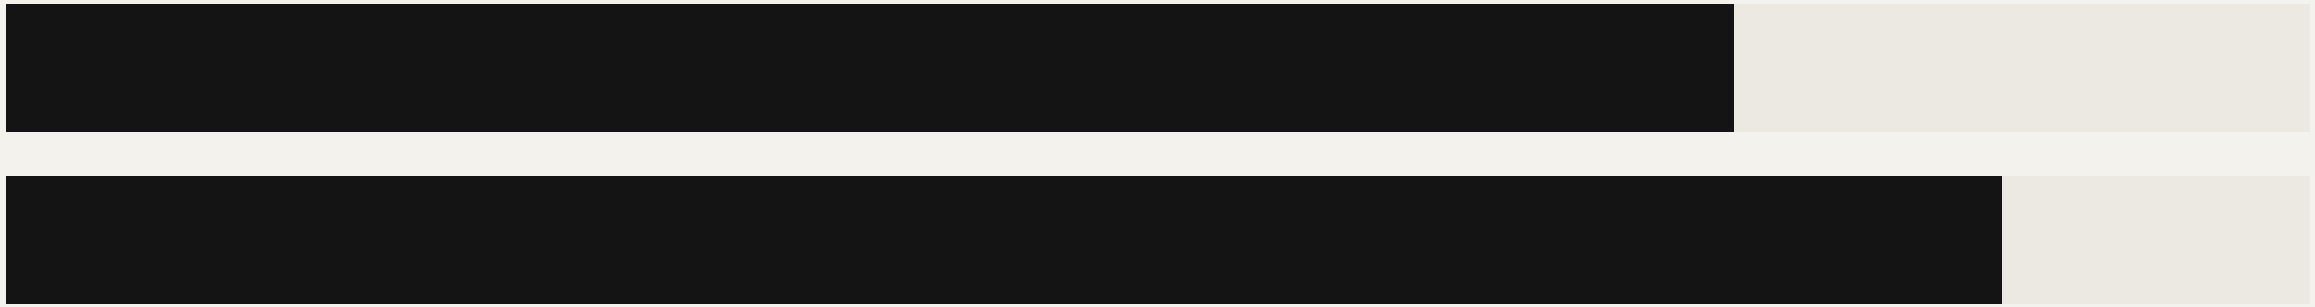
Look at the counts and the source before comparing the three runs.

They can agree perfectly and all recommend the wrong action.

Which score would catch that?

Which version would you ship?

First you graded your agent on six known activation questions. Now compare two supplied builds on a different set of eight requests. A gets six routes right and B gets seven. Read what B actually told the requester on E05, E06 and E07. Would you let a planner use that wording?



When good reviewers disagree

The board editor asks for a footnote crediting R17.

If a factual replacement meets the request, refuse the causal wording and offer the observed change.

If the board needs a causal conclusion, send it to Product Data for exposure evidence.

The eval owner records which job this case tests before changing the key.

Would this score survive a blind rerun?

Three agents can agree and all be wrong. Keep the approved key with the reviewer, compare each full answer with the key, and record how many runs agreed separately. Before signing, check what the agent retrieved while it was answering.